Sample PDF - page 1

The quick brown fox jumps over the lazy dog. This is a short sentence used to verify that text extraction and speech synthesis work end to end. Every word in this paragraph should be read aloud by the text to speech engine, producing a natural sounding sentence that a listener can follow without any visual reference to the original document.

Sample PDF - page 2

An audiobook is simply a recording of a written work being read aloud. With modern speech synthesis, it is possible to generate such recordings automatically from any text file or document. The quality depends on the voice model and the clarity of the source text, so well formatted documents generally produce better results than noisy scans.

Sample PDF - page 3

This application extracts text from a PDF document, splits it into paragraph-sized chunks, and feeds each chunk to a local neural text to speech engine. The resulting audio files are stored in a library that remembers your listening position between sessions, so you can always return to exactly the sentence where you stopped listening last time.

You can pause, resume, skip chapters, and come back later to continue exactly where you left off. The library is stored locally in a SQLite database, so no account or cloud service is required. Everything runs on your own computer, keeping your documents private and offline.

Sample PDF - page 5

Chapter based navigation is useful for long documents such as books and manuals. Each paragraph chunk becomes a separate chapter with its own duration, and the player shows both the elapsed time and the remaining time for the current chapter while it is being spoken aloud.

Speech synthesis engines have improved dramatically in recent years. Neural voice models now produce recordings that are clear, expressive and pleasant to listen to for extended periods, making them a practical alternative to professional narration for personal use cases.

Sample PDF - page 7

If the source PDF contains tables, footnotes or unusual layouts, the extracted text may contain artifacts. The converter applies basic cleanup such as removing standalone page numbers and joining wrapped lines, which handles the majority of everyday documents.

Sample PDF - page 8

This project is open source and built from freely available components: PyMuPDF for document parsing, Piper for neural speech synthesis, pygame for audio playback and PyQt6 for the desktop interface. Contributions and new voice models are always welcome.